

Mining Large Scale Datasets

Locality-Sensitive Hashing

(Adapted from CS246@Stanford.edu; <http://www.mmms.org>)

Sérgio Matos - aleixomatos@ua.pt


Locality-Sensitive Hashing

- Family of related techniques
- Allows to only examine pairs that are likely to be similar
 - Avoids quadratic growth in computation time

LSH: general idea

- Hash items into buckets using many different hash functions
 - ↪ Functions are designed so that similar items are more likely to hash into same bucket
- Only pairs that share a bucket for at least one of the hash functions need to be examined
 - ↪ There may be false negatives – pairs of similar items may not be considered at all
 - ↪ There may be false positives – pairs of items may be erroneously found as similar

Similar Documents: Steps



- **Shingling:** Converts a document into a set representation
- **Min-Hashing:** Convert large sets to short signatures, while preserving similarity
- **Locality-Sensitive Hashing:** Focus on pairs of signatures likely to be from similar documents

Min-Hashing

Permutations

2	4	3
3	2	4
7	1	7
6	3	2
1	6	6
5	7	1
4	5	5

Characteristic matrix

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0


Signature matrix

2	1	2	1
2	1	4	1
1	2	1	2


Similarities

1-2	1-3	2-4	3-4	
0	.75	.75	0	Original
0	.67	1	0	Signatures

Locality-Sensitive Hashing




Locality-Sensitive Hashing



LSH with a single min hash function (one band of one row)

Locality-Sensitive Hashing




LSH – Optimal scenario: only pairs of sets with similarity $> t$ are selected as candidates

Locality-Sensitive Hashing


- b bands, r rows/band
- Consider columns C1 and C2 with similarity s
- For any band (r rows):
 - Probability that all rows in band are equal $= s^r$
 - Probability that some row in band is unequal $= 1 - s^r$
 - Probability that no band identical $= (1 - s^r)^b$
 - Probability that at least 1 band identical $= 1 - (1 - s^r)^b$

Locality-Sensitive Hashing: S-curve



LSH – b bands of r rows: S-curve

Locality-Sensitive Hashing: S-curve



Given a fixed threshold t , we want to choose r and b such that $\text{Prob}(\text{Candidate pair})$ has a “step” around t .

LSH Families of Hash Functions

- In the context of LSH families, a “hash function” is any function that allows us to say whether two elements are candidates for comparison

- We use the notation

$$h(x) = h(y)$$

to mean “ h says x and y are a candidate pair”

- A family of hash functions is any set of hash functions from which we can pick one at random efficiently
 - For Min-Hashing signatures, each permutation of rows gives us a different Min-Hash function

LSH Families of Hash Functions


- Consider a space S of points with a distance measure $d(x, y)$
Can be Jaccard, Cosine, Euclidean, or other distance

(d_1, d_2, p_1, p_2) -sensitive family

- A family H of hash functions is said to be (d_1, d_2, p_1, p_2) -sensitive if for any x and y in S :
 - If $d(x, y) < d_1$, then the probability that $h(x) = h(y)$, over all $h \in H$, is at least p_1
 - If $d(x, y) > d_2$, then the probability that $h(x) = h(y)$, over all $h \in H$, is at most p_2

↪ With a LS Family we can do LSH!

(d_1, d_2, p_1, p_2) -sensitive family



For distances d_1 and below, the probability is at least p_1 , and for distances d_2 and above, the probability is at most p_2 . Between distances d_1 and d_2 , we know nothing.

Goal: Minimize difference btw d_1 and d_2 and maximize distance btw p_1 and p_2 .

Example of LS Family: Min-Hash

- Consider

S = space of all sets

d = Jaccard distance

H : a family of Min-Hash functions for all permutations of rows

- Then for any hash function $h \in H$

$$\Pr[h(x) = h(y)] = 1 - d(x, y)$$

Example of LS Family: Min-Hash

- For Jaccard distance, Min-Hashing gives a $(d_1, d_2, (1 - d_1), (1 - d_2))$ -sensitive family for any $d_1 < d_2$

Example:

H is a $(\underline{1/3}, 2/3, \underline{2/3}, 1/3)$ -sensitive family for S and d

If distance $\leq 1/3$, similarity $\geq 2/3$

Then probability that min-hash values
are the same is $\geq 2/3$

Amplifying a LS-Family

- Reproduce the “S-curve” effect for any LS family
- The “bands” technique we learned for signature matrices carries over to this more general setting
- Two constructions
 - AND** construction ~ “rows in a band”
 - OR** construction ~ “many bands”

AND of Hash Functions

- Given family H , construct family H' consisting of r functions from H

For $h = [h_1, \dots, h_r]$ in H' ,

$h(x) = h(y)$ if and only if $h_i(x) = h_i(y)$ for all i

\hookrightarrow corresponds to creating a band of size r

Theorem

If H is (d_1, d_2, p_1, p_2) -sensitive, H' is (d_1, d_2, p_1^r, p_2^r) -sensitive

Lowers probability for large distances (Good)

Also lowers probability for small distances (Bad)

OR of Hash Functions

- Given family H , construct family H' consisting of b functions from H

For $h = [h_1, \dots, h_b]$ in H' ,

$h(x) = h(y)$ if and only if $h_i(x) = h_i(y)$ for at least one i

Theorem

If H is (d_1, d_2, p_1, p_2) -sensitive,



H' is $(d_1, d_2, 1 - (1 - p_1)^b, 1 - (1 - p_2)^b)$ -sensitive

Raises probability for small distances (Good))

Raises probability for large distances (Bad)

Effect of AND and OR constructions

- **AND** makes all probabilities shrink
 - by choosing r correctly, we can make the lower probability approach 0 while the higher does not
- **OR** makes all probabilities grow
 - by choosing b correctly, we can make the upper probability approach 1 while the lower does not



Combining AND and OR Constructions

- r -way **AND** followed by b -way **OR** construction
- Same as in Min-Hashing
 - **AND**: If bands match in **all** r (values hash to same bucket)
 - **OR**: Columns that have **at least one** common bucket form a candidate pair
- If for points x and y $Pr[h(x) = h(y)] = s$
 - H will make (x, y) a candidate pair with probability s
 - r ***AND** b ***OR** construction makes (x, y) a candidate pair with probability $1 - (1 - s^r)^b$
 - ↪ **S-curve**
- Can use OR followed by AND; can combine sequences

AND-OR construction: example

s	$p = 1 - (1 - s^4)^4$
0.2	0.0064
0.3	0.0320
0.4	0.0985
0.5	0.2275
0.6	0.4260
0.7	0.6666
0.8	0.8785
0.9	0.9860

$r = 4$, $b = 4$ transforms a $(d_1, d_2, 0.8, 0.2)$ -sensitive family into a $(d_1, d_2, 0.8785, 0.0064)$ -sensitive family.

OR-AND construction: example

s	$p = (1 - (1 - s)^4)^4$
0.1	0.0140
0.2	0.1215
0.3	0.3334
0.4	0.5740
0.5	0.7725
0.6	0.9015
0.7	0.9680
0.8	0.9936


$r = 4$, $b = 4$ transforms a $(d_1, d_2, 0.8, 0.2)$ -sensitive family into a $(d_1, d_2, 0.9936, 0.1215)$ -sensitive family.

Cascading constructions

- Example: Apply the (4,4) OR-AND construction followed by the (4,4) AND-OR construction
- Transforms a $(d_1, d_2, 0.8, 0.2)$ -sensitive family into a $(d_1, d_2, 0.9999996, 0.0008715)$ -sensitive family

Note that this family uses 256 ($=4*4*4*4$) of the original hash functions

Constructions: visualization of threshold t

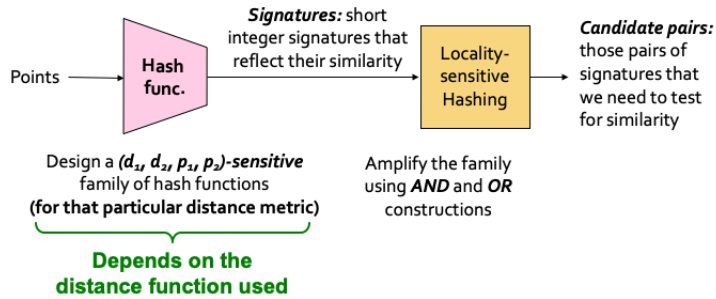


For an AND-OR S-curve $1 - (1 - s^r)^b$, the threshold t is where $1 - (1 - s^r)^b = t$
Probabilities p_1 and p_2 should be at opposite sides of t

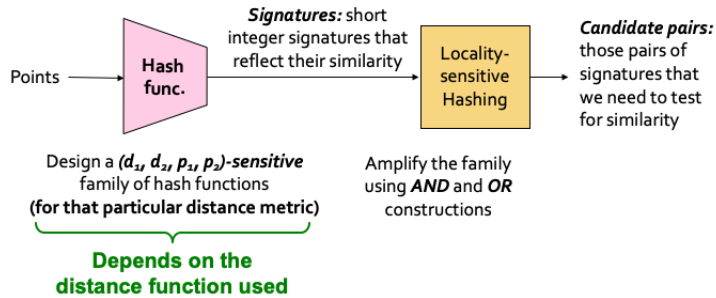
Summary

- Pick any two distances $d_1 < d_2$
- Start with a (d_1, d_2, p_1, p_2) -sensitive family
- Apply constructions to amplify (d_1, d_2, p_1^*, p_2^*) -sensitive family, where p_1^* is almost 1 and p_2^* is almost 0
- The closer to 0 and 1 we want to get, the more hash functions must be used!

LSH for other distance metrics



LSH for other distance metrics



- Cosine distance
 - **Random hyperplanes**
- Euclidean distance
 - **Random projections**

LSH for cosine distance

- For cosine distance, there is a technique called **Random Hyperplanes**
 - Technique similar to Min-Hashing
- Random Hyperplanes method is a $(d_1, d_2, (180 - d_1)/180, (180 - d_2)/180)$ - sensitive family for any d_1 and d_2


Reminder: (d_1, d_2, p_1, p_2) -sensitive

If $d(x, y) < d_1$, then probability that $h(x) = h(y)$ is at least p_1

If $d(x, y) > d_2$, then probability that $h(x) = h(y)$ is at most p_2


Random hyperplanes visualization

Look in the
plane of x
and y .



Note: what is important is that
hyperplane is outside the angle,
not that the vector is inside.

Random hyperplanes visualization



Signatures for Cosine Distance

- Pick some number of random vectors v_i , and hash your data for each vector

$$h_{v_i}(x) = +1 \text{ if } v_i \cdot x \geq 0$$

$$h_{v_i}(x) = -1 \text{ if } v_i \cdot x < 0$$


- Result is a signature (sketch) of +1's and -1's for each data point
- Can be used for LSH, in same way as Min-Hash signatures are used for Jaccard distance
- Amplify using AND/OR constructions

How to select random vectors


- Expensive to pick random vectors in M dimensions for large M
Would require generating M random numbers
- It suffices to consider only vectors v_i consisting of $+1$ and -1 components
Assuming data is random, then vectors of $+/-1$ cover the entire space evenly (no bias)

LSH for Euclidean distance


- Hash functions correspond to lines
- Partition the line into buckets (line segments) of size a
- Hash each point to the bucket containing its projection onto the line
- An element of the “signature” is a bucket id for that given projection line
- Nearby points are always close; distant points are rarely in same bucket



Projections



Projections




If $d \ll a$, then
the chance the
points are in the
same bucket is
at least $1 - d/a$.



Projections

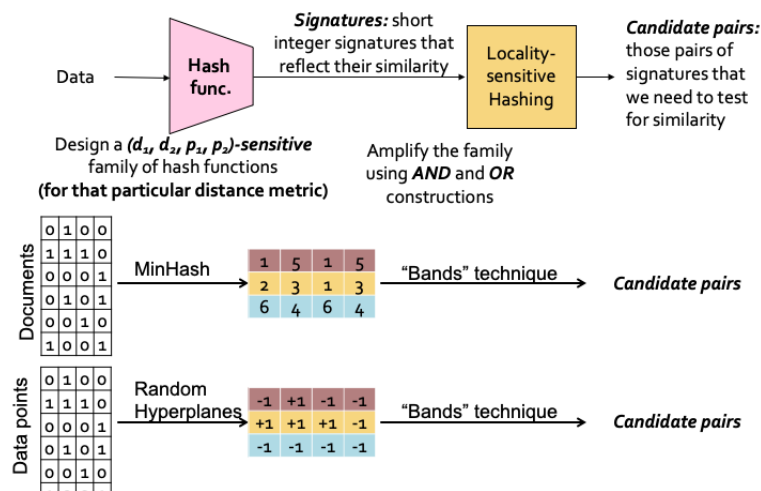
If $d \gg a$, θ must be close to 90° for there to be any chance points go to the same bucket.



LS-Family for Euclidean Distance

- If points are at distance $d \leq a/2$, then the probability of falling in same bucket is at least $1 - d/a$, or $1/2$
- If points are at distance $d \geq 2a$, then they can fall in the same bucket only if $d \cdot \cos\theta \leq a$, which is only true if
$$\cos\theta \leq 1/2$$
$$60^\circ < \theta < 90^\circ, \text{ i.e., at most } 1/3 \text{ probability}$$
- Yields a $(a/2, 2a, 1/2, 1/3)$ -sensitive family of hash functions for any a
- Amplify using AND-OR cascades

LS families



Important points

- Property $P(h(C1)=h(C2)) = \text{sim}(C1,C2)$ of hash function h is the essential part of LSH, without it we can't do anything
- LS-hash functions transform data to signatures so that the bands technique (AND, OR constructions) can then be applied

Further fun... [MMDS 3.8, 3.9]

- Application examples

- Entity resolution

- Idea: Score 100 for each full match on name, address, or phone

- Sort on name, score matching lines; repeat for address, phone

- Fingerprint matching

- Fingerprints represented by set of grid squares containing

- minutiae: can use Jaccard similarity

- LSH family: each hash function is a set of grid squares

- Similar news articles

- Shingles defined by a stopword followed by the next two words

- Sets as strings

- Length-based filtering

- Prefix indexing