



Mining Large Scale Datasets

Recommender Systems

(Adapted from CS246@Stanford.edu; <http://www.mmids.org>)

Sérgio Matos - aleixomatos@ua.pt

Recommendations



Examples:

amazon.com.



StumbleUpon



Google
News

last.fm
the social music revolution

XBOX
LIVE

You Tube

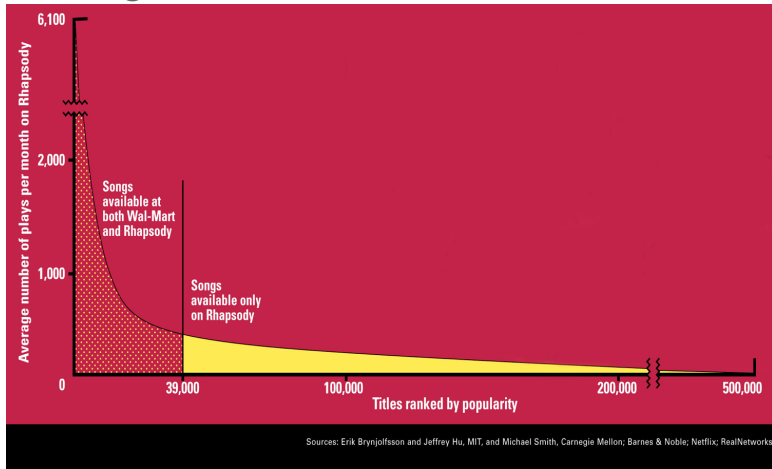
Scarcity vs Abundance

- Shelf space in traditional stores is scarce (and expensive)
 - Also: TV schedule, movie theaters, newspaper pages, ...
- Web enables near-zero-cost dissemination of information about products
 - ↪ Abundance


⇒ More choice necessitates better filters

- Recommendation engines
- Association rules:
 - How *Into Thin Air* made *Touching the Void* a bestseller

The long tail



The long tail




Physical vs Online

THE BIT PLAYER ADVANTAGE

Beyond bricks and mortar there are two main retail models – one that gets halfway down the Long Tail and another that goes all the way. The first is the familiar hybrid model of Amazon and Netflix, companies that sell physical goods online. Digital catalogs allow them to offer unlimited selection along with search, reviews, and recommendations, while the cost savings of massive warehouses and no walk-in customers greatly expands the number of products they can sell profitably.

Pushing this even further are pure digital services, such as iTunes, which offer the additional savings of delivering their digital goods online at virtually no marginal cost. Since an extra database entry and a few megabytes of storage on a server cost effectively nothing, these retailers have no economic reason not to carry everything available.



"IF YOU LIKE BRITNEY, YOU'LL LOVE ..."

Just as lower prices can entice consumers down the Long Tail, recommendation engines drive them to obscure content they might not find otherwise.



Types of recommendations

- Editorial and hand curated
 - List of favorites
 - Lists of “essential” items
- Simple aggregates
 - Top 10
 - Most Popular
 - Recent Uploads

⇒ **Tailored to individual users**

- Amazon, Netflix, ...

Formal model

- **X** = set of Customers
- **S** = set of Items
- Utility function $u : X \times S \rightarrow R$
 - R = set of ratings
 - R is a totally ordered set
 - e.g., 0-5 stars, real number in $[0,1]$

Utility matrix

	Avatar	LotR	Matrix	PotC
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Key problems

- (1) Gathering “known” ratings for matrix
 - How to collect the data in the utility matrix
- (2) Extrapolate unknown ratings from the known ones
 - Mainly interested in high unknown ratings
 - Interested in knowing what users like, not what they don't like
- (3) Evaluating extrapolation methods
 - How to measure success/performance of recommendation methods

Gathering ratings

- Explicit
 - Ask people to rate items
 - Doesn't work well in practice – most people won't be bothered; biased to those willing to rate
 - Crowdsourcing: Pay people to label items
- Implicit
 - Learn ratings from user actions
E.g., purchase / watching implies high rating
 - What about low ratings?

Extrapolating ratings

- Key problem: Utility matrix **U is sparse**
 - Most people have not rated most items
 - Cold start
 - New items have no ratings
 - New users have no history
- Three approaches to recommender systems
 - Content-based
 - Collaborative
 - Latent factor based

Content-based recommendation

- **Main idea**

Recommend to customer x items similar to previous items rated highly by x

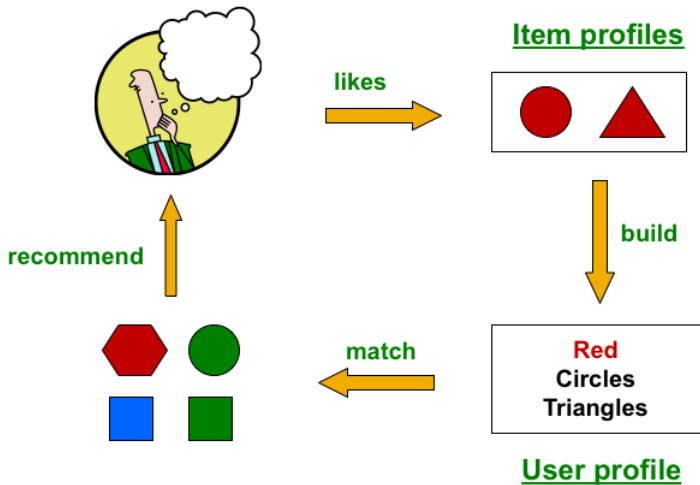
- Movie recommendations

Recommend movies with same actor(s), director, genre, ...

- Websites, blogs, news

Recommend other sites with “similar” content

Overview



Item profiles

- Create an **item profile** for each item
 - A set (vector) of features
 - Movies: author, title, actor, director, ...
 - Text: Set of “important” words in document
- How to pick important features?
 - Usual heuristic from text mining is TF-IDF
(Term frequency * Inverse Doc Frequency)
 - Doc profile = set of words with highest TF-IDF scores

User profiles and prediction

- User profile possibilities
 - Weighted average of rated item profiles
 - Variation: weight by difference from average rating for item
- Prediction heuristic: Cosine similarity of user and item profiles
Given user profile \mathbf{x} and item profile \mathbf{i} , estimate

$$u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{\|\mathbf{x}\| \cdot \|\mathbf{i}\|}$$

- How do you quickly find items closest to \mathbf{x} ?

User profiles and prediction

- User profile possibilities
 - Weighted average of rated item profiles
 - Variation: weight by difference from average rating for item
- Prediction heuristic: Cosine similarity of user and item profiles

Given user profile x and item profile i , estimate

$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{\|x\| \cdot \|i\|}$$

- How do you quickly find items closest to x ?
↪ LSH!

Content-based: Pros


- + No need for data on other users
- + Able to recommend to users with unique tastes
- + Able to recommend new and unpopular items
 - No first-rater problem
- + Able to provide explanations
 - Explain recommended items by listing content-features that caused items to be recommended

Content-based: Cons

- Finding the appropriate features is hard
 - E.g., images, movies, music
- Recommendations for new users
 - How to build a user profile?
- Overspecialization
 - Never recommends items outside user's content profile
 - People might have multiple interests
 - Unable to exploit quality judgments of other users

Collaborative filtering

- Consider user x
- Find set N of other users whose ratings are “similar” to x ’s ratings
- Estimate x ’s ratings based on ratings of the N users



Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Intuitively we want $\text{sim}(A, B) > \text{sim}(A, C)$

Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Jaccard similarity

$$\text{sim}(A, B) = 1/5 < 2/4 = \text{sim}(A, C)$$

- Problem: Ignores the values of ratings

Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Cosine similarity $sim(\mathbf{x}, \mathbf{y}) = \cos(\mathbf{r}_x, \mathbf{r}_y) = \frac{\mathbf{r}_x \cdot \mathbf{r}_y}{\|\mathbf{r}_x\| \cdot \|\mathbf{r}_y\|}$

$$sim(A, B) = 0.380 > 0.322 = sim(A, C)$$

- Problem: Treats missing ratings as “negative” (disliked)

$$r_A = 4, 0, 0, 5, 1, 0, 0, \quad r_B = 5, 5, 4, 0, 0, 0, 0$$

Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

- Cosine similarity $\text{sim}(\mathbf{x}, \mathbf{y}) = \cos(\mathbf{r}_x, \mathbf{r}_y) = \frac{\mathbf{r}_x \cdot \mathbf{r}_y}{\|\mathbf{r}_x\| \cdot \|\mathbf{r}_y\|}$
 $\text{sim}(A, B) = 0.380 > 0.322 = \text{sim}(A, C)$
- Problem: Treats missing ratings as “negative” (disliked)
- **Solution: subtract the (row) mean**
= Pearson correlation coefficient

Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Pearson correlation coefficient
 - S_{xy} = items rated by both users x and y

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$$

$$\text{sim}(A, B) = 0.092 > -0.559 = \text{sim}(A, C)$$

Predicting ratings

From similarity metric to recommendations

- Let \mathbf{r}_x be the vector of ratings for user x
- Let N be the set of k users most similar to x who have rated item i
- Prediction for item i of user x :

$$r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$$

or even better,

$$r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}} \quad , \quad s_{xy} = \text{sim}(x, y)$$

Item-Item Collaborative Filtering

- Item-item vs User-user
- For item i , find other similar items
- Estimate rating for item i based on ratings for similar items
- Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

s_{ij} : similarity of items i and j

r_{xj} : rating of user x on item j

$N(i;x)$: set items rated by x that are similar to i

Item-Item CF

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3			5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- unknown rating



- rating between 1 to 5

Item-Item CF

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

Item-Item CF

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	$\text{sim}(1,m)$
movies	1	1		3		?	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		<u>0.59</u>

Neighbor selection:

Identify movies similar to
movie 1, rated by user 5

Here we use Pearson correlation as similarity:

1) Subtract mean rating m_i from each movie i

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]

2) Compute cosine similarities between rows

Item-Item CF

		users													
		1	2	3	4	5	6	7	8	9	10	11	12	$\text{sim}(1,m)$	
movies	1	1		3		?	5			5		4		1.00	
	2			5	4			4			2	1	3	-0.18	
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>	
	4		2	4		5			4			2		-0.10	
	5			4	3	4	2					2	5	-0.31	
	<u>6</u>	1		3		3			2			4		<u>0.59</u>	

Compute similarity weights:

$$s_{1,3}=0.41, s_{1,6}=0.59$$

Item-Item CF

movies	users											
	1	2	3	4	5	6	7	8	9	10	11	12
	1		3		2.6	5			5		4	
	2		5	4			4			2	1	3
	<u>3</u>	2	4		1	2		3		4	3	5
	4		2	4		5			4		2	
	5			4	3	4	2				2	5
	<u>6</u>	1		3		3			2		4	

Predict by taking weighted average:

$$r_{1.5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

$$r_{ix} = \frac{\sum_{j \in N(i,x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$

Item-item vs User-user

	Avatar	LotR	Matrix	PotC
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

- In theory, these are dual approaches with similar performance
- In practice, it has been observed that item-item often works better than user-user
- Why? Items are simpler, users have multiple tastes

Pros/Cons of Collaborative Filtering

- + Works for any kind of item
 - No feature selection needed
- Cold Start
 - Need enough users in the system to find a match
- Sparsity
 - The user/ratings matrix is sparse
 - Hard to find users that have rated the same items
- First rater
 - Cannot recommend an item that has not been previously rated
 - New items, esoteric items
- Popularity bias
 - Cannot recommend items to someone with unique taste
 - Tends to recommend popular items

Hybrid methods

Combine predictions from two or more different recommenders

- e.g. Global baseline + CF
- Perhaps using a linear model

Add content-based methods to collaborative filtering

- Item profiles for new item problem
- Demographics to deal with new user problem

CF: Common practice

- Define similarity s_{ij} of items i and j
- Select k nearest neighbors $N(i; x)$
 - Items most similar to i , that were rated by x
- Estimate rating r_{xi} as the weighted average

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$

$b_{xi} = \mu + b_x + b_i$ baseline estimate for r_{xi}

μ = overall mean movie rating

b_x = rating deviation of user x = (avg rating of user x) - μ

b_i = rating deviation of movie i = (avg rating of movie i) - μ

Evaluation

A 10x6 grid representing user-movie ratings. The grid is labeled "users" on the left and "movies" on the top. The grid contains numerical ratings from 1 to 5 in some cells, while others are empty.

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

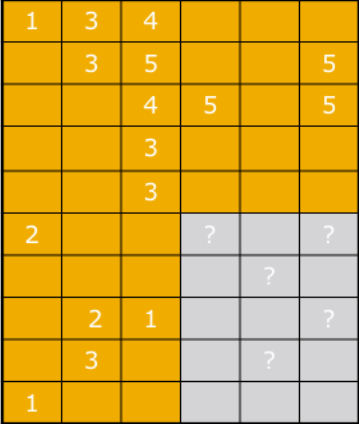
Evaluation

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

Test Data Set



Evaluating predictions

- Compare predictions with known ratings

- Root-mean-square error (RMSE)

$$\sqrt{\frac{\sum_{xi} (r_{xi} - r_{xi}^*)^2}{\sum_{xi} 1}}$$

where r_{xi} is predicted; r_{xi}^* is the true rating

- Precision at top 10
 - Rank correlation

Spearman's correlation between system's and user's complete rankings

- Another approach: 0/1 model (dislike/like)

- Coverage

items/users for which the system can make predictions

- Precision
 - Receiver operating characteristic (ROC)

Tradeoff curve between false positives and false negatives


Problems with error measures

- Narrow focus on accuracy sometimes misses the point
 - Prediction diversity
 - Prediction context
 - Order of predictions
- In practice, we care only about predicting high ratings
 - RMSE might penalize a method that does well for high ratings and badly for others

Collaborative Filtering: Complexity

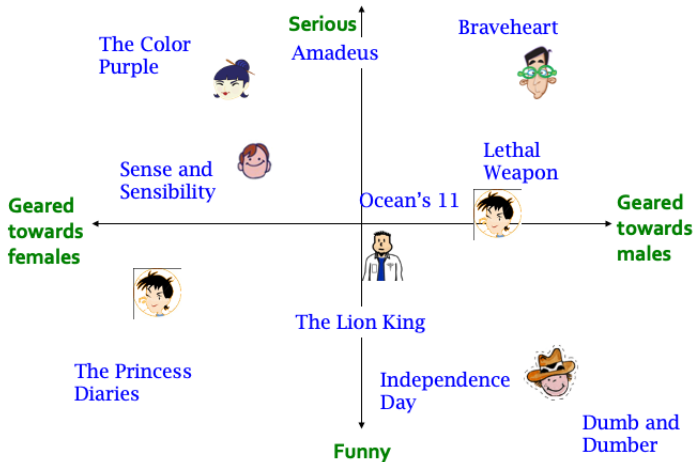
- Expensive step is finding k most similar customers: $O(|X|)$
- Too expensive to do at runtime
 - Could pre-compute
 - Naïve pre-computation takes time $O(k \cdot |X|)$
- We already know how to do this!
 - Near-neighbor search in high dimensions (LSH)
 - Clustering
 - Dimensionality reduction (PCA, SVD)

Latent factor models




Assume we can approximate the rating matrix R as a product of "thin" $Q \cdot P^T$


Latent factor models




Latent factor models



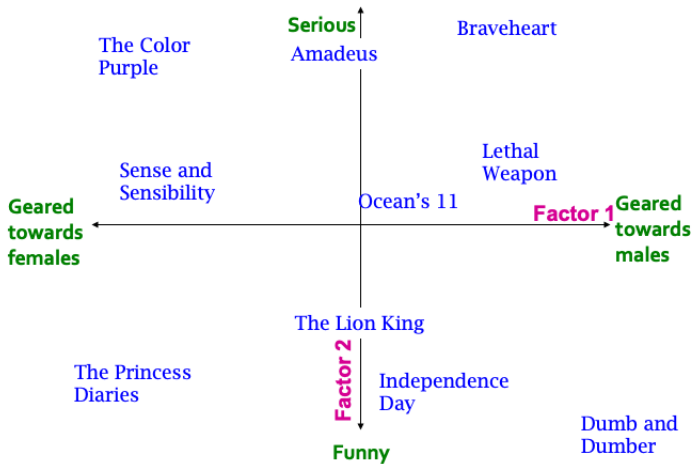
Latent factor models



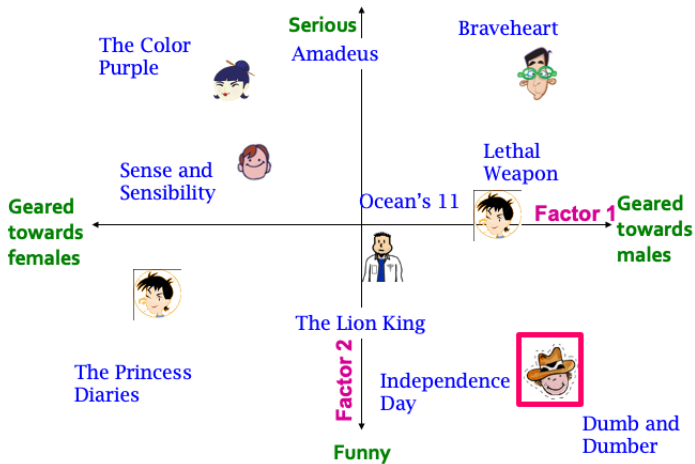
Latent factor models




Latent factor models



Latent factor models




Latent factor models: SVD



- A : Input matrix
- U : Left singular matrix
- V : Right singular matrix
- Σ : Singular values

Latent factors model: $R \approx Q \cdot P^T$
As SVD: $A = R, Q = U, P^T = \Sigma V^T$

Latent factor models: SVD




- A : Input matrix
- U : Left singular matrix
- V : Right singular matrix
- Σ : Singular values

Latent factors model: $R \approx Q \cdot P^T$

As SVD: $A = R, Q = U, P^T = \Sigma V^T$

\hookrightarrow SVD minimizes SSE, and thus RMSE!


Latent factor models: SVD



X SVD is not defined when entries are missing!

- Need a special method to find P , Q

Latent factor models



X SVD is not defined when entries are missing!

- Need a special method to find P, Q

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2$$

- We don't require cols of P, Q to be orthogonal/unit length
- P, Q map users/movies to a latent space

Latent factor models

	users										
items	1		3			5			5		4
			5	4			4			2	1
	2	4		1	2		3		4	3	5
		2	4		5			4			2
			4	3	4	2				2	5
	1		3		3			2			4
	R										

\approx

	factors			
items	.1	-.4	.2	
	-.5	.6	.5	
	-.2	.3	.5	
	1.1	2.1	.3	
	-.7	2.1	-2	
	-1	.7	.3	
	Q			

	users										
factors	1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4
	-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1
	2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6
	P^T										

Goal is to find P, Q such that $\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2$

Want to minimize SSE for unseen test data

- Approach: Minimize SSE on training data
- Want large k (number of factors) to capture all the signals
- But SSE on test data begins to rise for $k > 2$
 \hookrightarrow **Overfitting**

Latent factor models

Solution: **Regularization**

- Allow rich model where there are sufficient data
- Shrink aggressively where data are scarce

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2 + \left[\lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 \right]$$

Latent factor models

Combining with baseline predictor

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

- μ : overall mean rating
- b_x : bias of user x
- b_i : bias of movie i

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - (\mu + b_x + b_i + q_i \cdot p_x))^2 +$$
$$\left[\lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 + \lambda_3 \sum_x \|b_x\|^2 + \lambda_4 \sum_i \|b_i\|^2 \right]$$

Latent factor models

Combining with baseline predictor

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

- μ : overall mean rating
- b_x : bias of user x
- b_i : bias of movie i

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - (\mu + b_x + b_i + q_i \cdot p_x))^2 +$$
$$\left[\lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 + \lambda_3 \sum_x \|b_x\|^2 + \lambda_4 \sum_i \|b_i\|^2 \right]$$

↪ **Stochastic Gradient Descent**

Final tip: Add data

- Leverage all the data
 - Don't try to reduce data size in an effort to make fancy algorithms work
 - Simple methods on large data do best
- Add more data
 - e.g., add IMDB data on genres
- **More Richer data beats better algorithms**