


# Mining Large Scale Datasets

## Recommender Systems

(Adapted from CS246@Stanford.edu; <http://www.mmms.org>)

Sérgio Matos - [aleixomatos@ua.pt](mailto:aleixomatos@ua.pt)

# Recommendations



## Examples:

amazon.com.



StumbleUpon



Google  
News

last.fm  
the social music revolution

XBOX  
LIVE

You Tube


# Scarcity vs Abundance

- Shelf space in traditional stores is scarce (and expensive)
  - Also: TV schedule, movie theaters, newspaper pages, ...
- Web enables near-zero-cost dissemination of information about products
  - ↪ Abundance


⇒ More choice necessitates better filters

- Recommendation engines
- Association rules:  
How *Into Thin Air* made *Touching the Void* a bestseller

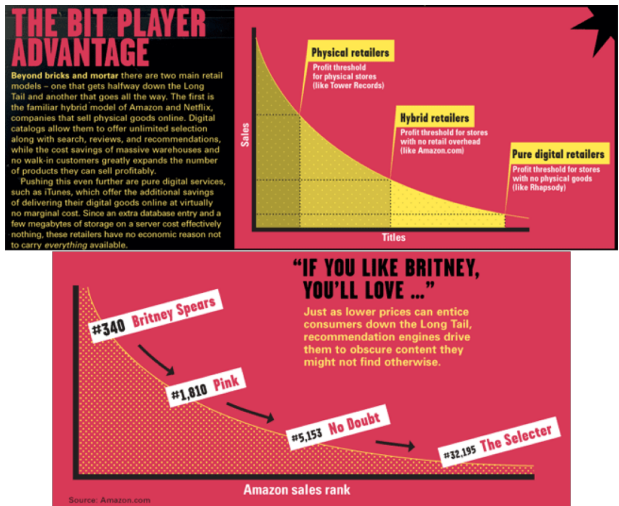
# The long tail



# The long tail



# Physical vs Online



# Types of recommendations

- Editorial and hand curated
  - List of favorites
  - Lists of “essential” items
- Simple aggregates
  - Top 10
  - Most Popular
  - Recent Uploads

⇒ **Tailored to individual users**

- Amazon, Netflix, ...

# Formal model

- **X** = set of Customers
- **S** = set of Items
- Utility function  $u : X \times S \rightarrow R$ 
  - R = set of ratings
  - R is a totally ordered set
  - e.g., 0-5 stars, real number in  $[0,1]$



# Utility matrix

	Avatar	LotR	Matrix	PotC
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

# Key problems

- (1) Gathering “known” ratings for matrix
  - How to collect the data in the utility matrix
- (2) Extrapolate unknown ratings from the known ones
  - Mainly interested in high unknown ratings
  - Interested in knowing what users like, not what they don't like
- (3) Evaluating extrapolation methods
  - How to measure success/performance of recommendation methods

# Gathering ratings

- Explicit
  - Ask people to rate items
  - Doesn't work well in practice – most people won't be bothered; biased to those willing to rate
  - Crowdsourcing: Pay people to label items
- Implicit
  - Learn ratings from user actions
    - E.g., purchase / watching implies high rating
  - What about low ratings?

# Extrapolating ratings

- Key problem: Utility matrix **U is sparse**
  - Most people have not rated most items
  - Cold start
    - New items have no ratings
    - New users have no history
- Three approaches to recommender systems
  - Content-based
  - Collaborative
  - Latent factor based

# Content-based recommendation

- **Main idea**

Recommend to customer  $x$  items similar to previous items rated highly by  $x$

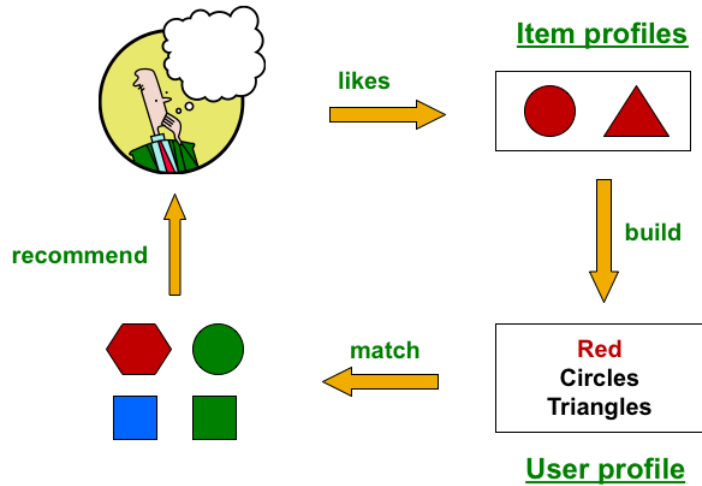
- Movie recommendations

Recommend movies with same actor(s), director, genre, ...

- Websites, blogs, news

Recommend other sites with “similar” content

# Overview



# Item profiles

- Create an **item profile** for each item
  - A set (vector) of features
  - Movies: author, title, actor, director, ...
  - Text: Set of “important” words in document
- How to pick important features?
  - Usual heuristic from text mining is TF-IDF (Term frequency \* Inverse Doc Frequency)
  - Doc profile = set of words with highest TF-IDF scores

# User profiles and prediction

- User profile possibilities
  - Weighted average of rated item profiles
  - Variation: weight by difference from average rating for item
- Prediction heuristic: Cosine similarity of user and item profiles  
Given user profile  $\mathbf{x}$  and item profile  $\mathbf{i}$ , estimate

$$u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{\|\mathbf{x}\| \cdot \|\mathbf{i}\|}$$

- How do you quickly find items closest to  $\mathbf{x}$ ?



# User profiles and prediction

- User profile possibilities
  - Weighted average of rated item profiles
  - Variation: weight by difference from average rating for item
- Prediction heuristic: Cosine similarity of user and item profiles

Given user profile  $\mathbf{x}$  and item profile  $\mathbf{i}$ , estimate

$$u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{\|\mathbf{x}\| \cdot \|\mathbf{i}\|}$$

- How do you quickly find items closest to  $\mathbf{x}$ ?  
↪ LSH!

## Content-based: Pros


- + No need for data on other users
- + Able to recommend to users with unique tastes
- + Able to recommend new and unpopular items
  - No first-rater problem
- + Able to provide explanations
  - Explain recommended items by listing content-features that caused items to be recommended

## Content-based: Cons

- Finding the appropriate features is hard
  - E.g., images, movies, music
- Recommendations for new users
  - How to build a user profile?
- Overspecialization
  - Never recommends items outside user's content profile
  - People might have multiple interests
  - Unable to exploit quality judgments of other users

# Collaborative filtering

- Consider user  $x$
- Find set  $N$  of other users whose ratings are “similar” to  $x$ ’s ratings
- Estimate  $x$ ’s ratings based on ratings of the  $N$  users



## Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- Intuitively we want  $\text{sim}(A, B) > \text{sim}(A, C)$

## Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Jaccard similarity

$$\text{sim}(A, B) = 1/5 < 2/4 = \text{sim}(A, C)$$

- Problem: Ignores the values of ratings

## Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Cosine similarity  $\text{sim}(\mathbf{x}, \mathbf{y}) = \cos(\mathbf{r}_x, \mathbf{r}_y) = \frac{\mathbf{r}_x \cdot \mathbf{r}_y}{\|\mathbf{r}_x\| \cdot \|\mathbf{r}_y\|}$   
 $\text{sim}(A, B) = 0.380 > 0.322 = \text{sim}(A, C)$
- Problem: Treats missing ratings as “negative” (disliked)  
 $r_A = 4, 0, 0, 5, 1, 0, 0, r_B = 5, 5, 4, 0, 0, 0, 0$

## Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

- Cosine similarity  $sim(\mathbf{x}, \mathbf{y}) = \cos(\mathbf{r}_x, \mathbf{r}_y) = \frac{\mathbf{r}_x \cdot \mathbf{r}_y}{\|\mathbf{r}_x\| \cdot \|\mathbf{r}_y\|}$   
 $sim(A, B) = 0.380 > 0.322 = sim(A, C)$
- Problem: Treats missing ratings as “negative” (disliked)
- **Solution: subtract the (row) mean**  
= Pearson correlation coefficient



## Finding “similar” users: Similarity metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Pearson correlation coefficient
  - $S_{xy}$  = items rated by both users  $x$  and  $y$

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$$

$$\text{sim}(A, B) = 0.092 > -0.559 = \text{sim}(A, C)$$

# Predicting ratings

From similarity metric to recommendations

- Let  $\mathbf{r}_x$  be the vector of ratings for user  $x$
- Let  $N$  be the set of  $k$  users most similar to  $x$  who have rated item  $i$
- Prediction for item  $i$  of user  $x$ :

$$r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$$

or even better,

$$r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}} \quad , \quad s_{xy} = \text{sim}(x, y)$$

# Item-Item Collaborative Filtering

- Item-item vs User-user
- For item  $i$ , find other similar items
- Estimate rating for item  $i$  based on ratings for similar items
- Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

$s_{ij}$ : similarity of items  $i$  and  $j$

$r_{xj}$ : rating of user  $x$  on item  $j$

$N(i;x)$ : set items rated by  $x$  that are similar to  $i$

# Item-Item CF

movies	users											
	1	2	3	4	5	6	7	8	9	10	11	12
	1		3			5			5		4	
	2		5	4			4			2	1	3
	3	2	4		1	2	3		4	3	5	
	4		2	4		5		4			2	
	5		4	3	4	2					2	5
6	1		3		3			2			4	



- unknown rating



- rating between 1 to 5

# Item-Item CF

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

# Item-Item CF

		users												sim(1,m)
		1	2	3	4	5	6	7	8	9	10	11	12	
movies	1	1		3		?	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	3	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	6	1		3		3			2			4		<u>0.59</u>

**Neighbor selection:**  
Identify movies similar to  
movie 1, rated by user 5

Here we use Pearson correlation as similarity:  
1) Subtract mean rating  $m_i$  from each movie  $i$   
 $m_1 = (1+3+5+5+4)/5 = 3.6$   
row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]  
2) Compute cosine similarities between rows

# Item-Item CF

	users												
	1	2	3	4	5	6	7	8	9	10	11	12	sim(1,m)
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	-0.18
<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
<u>6</u>	1		3		3			2			4		<u>0.59</u>

Compute similarity weights:

$s_{1,3}=0.41$ ,  $s_{1,6}=0.59$

# Item-Item CF

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		2.6	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	

Predict by taking weighted average:

$$r_{1,5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6 \quad r_{ix} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$



## Item-item vs User-user

	Avatar	LotR	Matrix	PotC
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

- In theory, these are dual approaches with similar performance
- In practice, it has been observed that item-item often works better than user-user
- Why? Items are simpler, users have multiple tastes

# Pros/Cons of Collaborative Filtering

- + Works for any kind of item
  - No feature selection needed
- Cold Start
  - Need enough users in the system to find a match
- Sparsity
  - The user/ratings matrix is sparse
  - Hard to find users that have rated the same items
- First rater
  - Cannot recommend an item that has not been previously rated
  - New items, esoteric items
- Popularity bias
  - Cannot recommend items to someone with unique taste
  - Tends to recommend popular items

# Hybrid methods

Combine predictions from two or more different recommenders

- e.g. Global baseline + CF
- Perhaps using a linear model

Add content-based methods to collaborative filtering

- Item profiles for new item problem
- Demographics to deal with new user problem

## CF: Common practice

- Define similarity  $s_{ij}$  of items  $i$  and  $j$
- Select  $k$  nearest neighbors  $N(i; x)$ 
  - Items most similar to  $i$ , that were rated by  $x$
- Estimate rating  $r_{xi}$  as the weighted average

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$

$b_{xi} = \mu + b_x + b_i$  baseline estimate for  $r_{xi}$

$\mu$  = overall mean movie rating

$b_x$  = rating deviation of user  $x$  = (avg rating of user  $x$ ) -  $\mu$

$b_i$  = rating deviation of movie  $i$  = (avg rating of movie  $i$ ) -  $\mu$

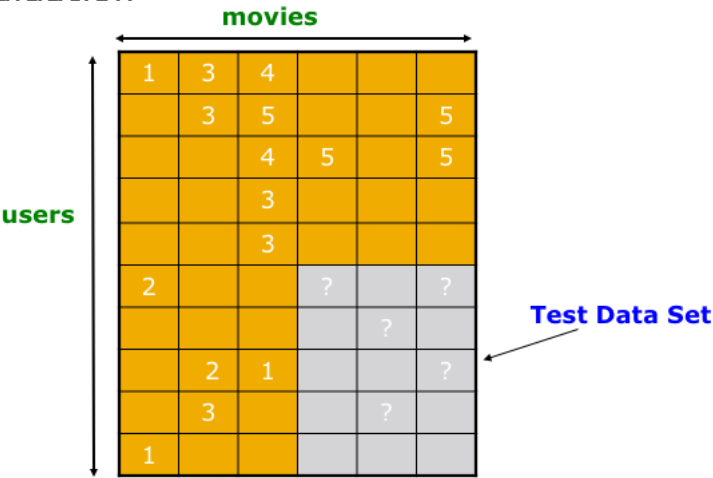
# Evaluation

movies

users

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

# Evaluation



# Evaluating predictions

- Compare predictions with known ratings

- Root-mean-square error (RMSE)

$$\sqrt{\frac{\sum_{xj} (r_{xj} - r_{xj}^*)^2}{\sum_{xj} 1}} \quad \text{where } r_{xj} \text{ is predicted; } r_{xj}^* \text{ is the true rating}$$

- Precision at top 10

- Rank correlation

Spearman's correlation between system's and user's complete rankings

- Another approach: 0/1 model (dislike/like)

- Coverage

# items/users for which the system can make predictions

- Precision

- Receiver operating characteristic (ROC)

Tradeoff curve between false positives and false negatives

# Problems with error measures


- Narrow focus on accuracy sometimes misses the point
  - Prediction diversity
  - Prediction context
  - Order of predictions
- In practice, we care only about predicting high ratings
  - RMSE might penalize a method that does well for high ratings and badly for others



## Collaborative Filtering: Complexity

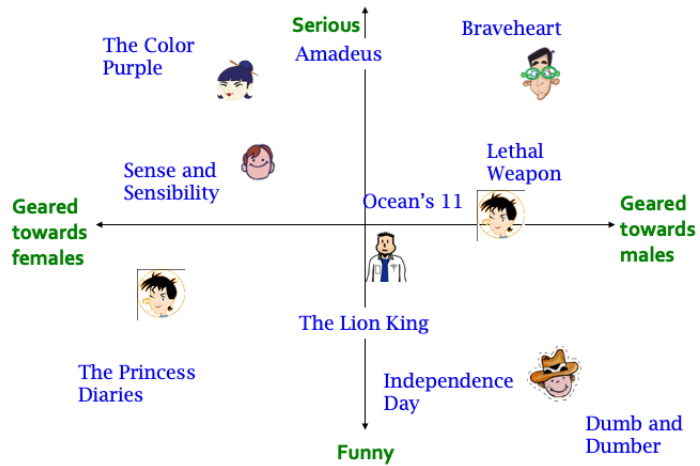
- Expensive step is finding  $k$  most similar customers:  $O(|X|)$
- Too expensive to do at runtime
  - Could pre-compute
  - Naïve pre-computation takes time  $O(k \cdot |X|)$
- We already know how to do this!
  - Near-neighbor search in high dimensions (LSH)
  - Clustering
  - Dimensionality reduction (PCA, SVD)

# Latent factor models




Assume we can approximate the rating matrix  $R$  as a product of “thin”  $Q \cdot P^T$


# Latent factor models




# Latent factor models



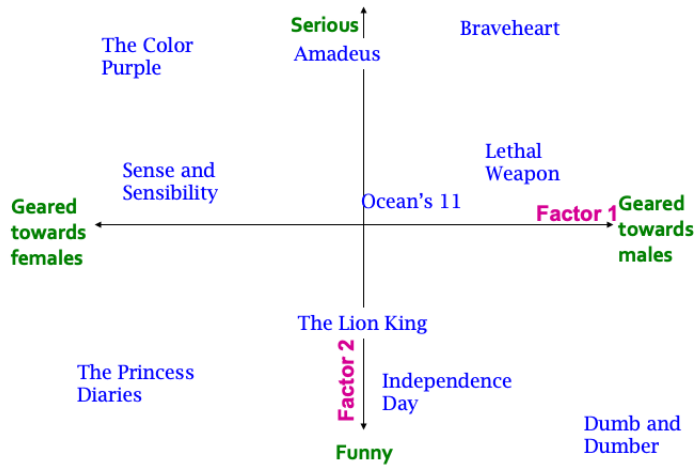
# Latent factor models



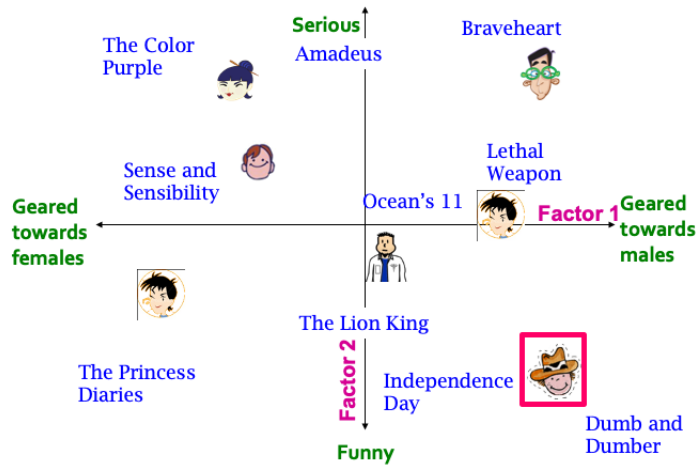
# Latent factor models



# Latent factor models




# Latent factor models






## Latent factor models: SVD



Latent factors model:  $R \approx Q \cdot P^T$

As SVD:  $A = R$ ,  $Q = U$ ,  $P^T = \Sigma V^T$

## Latent factor models: SVD




- $A$ : Input matrix
- $U$ : Left singular matrix
- $V$ : Right singular matrix
- $\Sigma$ : Singular values

Latent factors model:  $R \approx Q \cdot P^T$

As SVD:  $A = R$ ,  $Q = U$ ,  $P^T = \Sigma V^T$

↪ SVD minimizes SSE, and thus RMSE!


# Latent factor models: SVD



✗ SVD is not defined when entries are missing!

- Need a special method to find  $P$ ,  $Q$

# Latent factor models




X SVD is not defined when entries are missing!

- Need a special method to find  $P$ ,  $Q$

$$\min_{P, Q} \sum_{(i, x) \in R} (r_{xi} - q_i \cdot p_x)^2$$

- We don't require cols of  $P$ ,  $Q$  to be orthogonal/unit length
- $P$ ,  $Q$  map users/movies to a latent space

# Latent factor models



Goal is to find  $P, Q$  such that  $\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2$

Want to minimize SSE for unseen test data

- Approach: Minimize SSE on training data
- Want large  $k$  (number of factors) to capture all the signals
- But SSE on test data begins to rise for  $k > 2$

↪ **Overfitting**

# Latent factor models

Solution: **Regularization**

- Allow rich model where there are sufficient data
- Shrink aggressively where data are scarce

$$\min_{P,Q} \sum_{(i,x) \in R} (r_{xi} - q_i \cdot p_x)^2 + \left[ \lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 \right]$$

# Latent factor models

Combining with baseline predictor

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

- $\mu$ : overall mean rating
- $b_x$ : bias of user x
- $b_i$ : bias of movie i

$$\min_{P, Q} \sum_{(i,x) \in R} (r_{xi} - (\mu + b_x + b_i + q_i \cdot p_x))^2 +$$
$$\left[ \lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 + \lambda_3 \sum_x \|b_x\|^2 + \lambda_4 \sum_i \|b_i\|^2 \right]$$

# Latent factor models

Combining with baseline predictor

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

- $\mu$ : overall mean rating
- $b_x$ : bias of user  $x$
- $b_i$ : bias of movie  $i$

$$\min_{P, Q} \sum_{(i, x) \in R} (r_{xi} - (\mu + b_x + b_i + q_i \cdot p_x))^2 +$$
$$\left[ \lambda_1 \sum_x \|p_x\|^2 + \lambda_2 \sum_i \|q_i\|^2 + \lambda_3 \sum_x \|b_x\|^2 + \lambda_4 \sum_i \|b_i\|^2 \right]$$

→ **Stochastic Gradient Descent**



## Final tip: Add data

- Leverage all the data
  - Don't try to reduce data size in an effort to make fancy algorithms work
  - Simple methods on large data do best
- Add more data
  - e.g., add IMDB data on genres
- **More Richer data beats better algorithms**