


Mining Large Scale Datasets

Graph Representation Learning

(Adapted from CS246@Stanford.edu; <http://www.mmids.org>)


Sérgio Matos - aleixomatos@ua.pt

Machine Learning on Graphs



Node classification

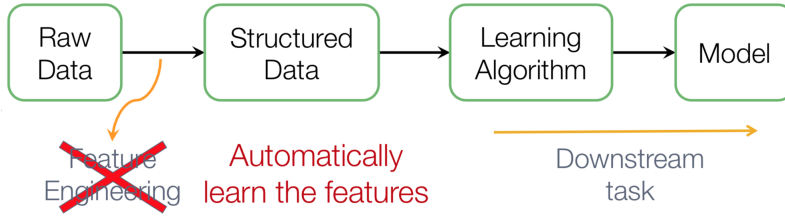
Machine Learning on Graphs



Link prediction


Machine Learning on Graphs

Machine Learning requires feature engineering!



Feature Learning on Graphs


GOAL: Efficient task-independent feature learning for machine learning with graphs




Feature Learning on Graphs

TASK: Map each node in a network into a low-dimensional space


- Distributed representation of nodes
- Similarity of embeddings (node representations) between nodes indicates their network similarity
- Encode network information and generate node representation



Feature Learning on Graphs: Example



(a) Input: Karate Graph




(b) Output: Representation

Image from: Perozzi et al. Deepwalk 2014


Feature Learning on Graphs: Differences

- Modern deep learning toolbox is designed for simple sequences or grids
 - CNNs for fixed-size images
 - RNNs or word2vec for text/sequences
- But networks are far more complex
 - Complex topographical structure: no spatial locality like grids
 - No fixed node ordering or reference point
 - Often dynamic and with multimodal features




Embedding graph nodes

- Assume a graph G
 - V is the vertex set
 - \mathbf{A} is the (binary) adjacency matrix
- Goal is to encode nodes so that similarity in the embedding space (e.g., dot product) approximates similarity in the original network



Embedding graph nodes

- Assume a graph G
 - V is the vertex set
 - \mathbf{A} is the (binary) adjacency matrix
- Goal is to encode nodes so that similarity in the embedding space (e.g., dot product) approximates similarity in the original network



Embedding graph nodes: Steps

- 1 Define an encoder (i.e., a mapping from nodes to embeddings)
- 2 Define a node similarity function (i.e., a measure of similarity in the original network)
- 3 Optimize the parameters of the encoder so that

$$\text{similarity}(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$$

Embedding graph nodes: Key components

- Encoder

- Maps a node v to a low-dimensional vector \mathbf{z}_v

$$ENC(v) = \mathbf{z}_v$$

- Similarity function

- Specifies how the relationships in vector space map to the relationships in the original network

$$similarity(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$$

Shallow encoding

- Simplest encoding approach: encoder is just an embedding-lookup


$$ENC(v) = \mathbf{Z} \cdot \mathbf{v}$$

$\mathbf{Z} \in \mathbb{R}^{d \times |\mathcal{V}|}$ matrix, each column is a node embedding
 \hookrightarrow what we learn!

$\mathbf{v} \in \mathbb{I}^{|\mathcal{V}|}$ indicator vector
all 0s except a 1 in column indicating node v


Shallow encoding

- Simplest encoding approach: encoder is just an embedding-lookup



Shallow encoding

- Simplest encoding approach: encoder is just an embedding-lookup




Methods: DeepWalk, node2vec, TransX, ...

Node similarity


- Key choice of methods is how they define node similarity
- E.g., should two nodes have similar embeddings if they ...
 - Are connected?
 - Share neighbours?
 - Have similar “structural roles”?
 - ...?

Random walk approaches



- Random walk on a graph
 - Given a starting point
 - Select a neighbor at random and move to that node
 - Repeat


Random walk approaches



- Random walk on a graph
 - Given a starting point
 - Select a neighbor at random and move to that node
 - Repeat
- $\mathbf{z}_u^T \mathbf{z}_v \approx$ probability that u and v co-occur on a random walk over the graph

Random walk embeddings


1. Estimate probability of visiting node v on a random walk starting from node u using some random walk strategy R



$$P_R(v|u)$$

2. Optimize embeddings to encode these random walk statistics

Note: $\mathbf{z}_u^\top \mathbf{z}_v = \cos(\theta)$ encodes the random walk 'similarity'



$$\theta \propto P_R(v|u)$$

Random walks

- **Expressivity**

Flexible stochastic definition of node similarity that incorporates both local and higher-order neighborhood information

- **Efficiency**

No need to consider all node pairs when training

Only need to consider pairs that co-occur on random walks

Unsupervised feature learning

Intuition: Find d -dimensional embedding of nodes that preserves similarity

Idea: Learn node embedding such that nearby nodes are close together in encoding space

Given a node u how do we define nearby nodes?

$N_R(u)$: neighbourhood of u obtained by some strategy R

Feature learning as Optimization

Given $G = (V, E)$

Our goal is to learn a mapping $z : u \rightarrow R^d$

Log-likelihood objective:

$$\max_z \sum_{u \in V} \log P(N_R(u) | z_u)$$

where $N_R(u)$ is the neighborhood of node u by strategy R

Given node u we want to learn feature representations that are predictive of the nodes in its neighborhood $N_R(u)$

Random Walk Optimization

1. Run short fixed-length random walks starting from each node on the graph using some strategy R
2. For each node u collect $N_R(u)$, the multiset of nodes visited on random walks starting from u

Note: $N_R(u)$ can have repeated elements since nodes can be visited multiple times on random walks

3. Optimize embeddings according to
(Given node u , predict its neighbors $N_R(u)$)

$$\max_z \sum_{u \in V} \log P(N_R(u) | z_u)$$

Random Walk Optimization

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$

- **Intuition**

Optimize embeddings to maximize likelihood of random walk occurrences

- Parameterize $P(v|\mathbf{z}_u)$ using softmax

$$P(v|\mathbf{z}_u) = \frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \mathbf{z}_u^\top \mathbf{z}_n}$$

Random Walk Optimization

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} - \log \left(\frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^\top \mathbf{z}_n)} \right)$$

sum over all nodes u sum over nodes v seen on random walks starting from u predicted probability of u and v co-occurring on random walk

Optimizing random walk embeddings = finding embeddings \mathbf{z}_u that minimize \mathcal{L}

Random Walk Optimization

But doing this naively is too expensive!!

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log \left(\frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^\top \mathbf{z}_n)} \right)$$

Nested sum over nodes means $O(|V|^2)$ complexity!

Negative Sampling

- **Solution:** Negative sampling

Instead of normalizing w.r.t. all nodes, just normalize against k random “negative samples” n_i

$$\log\left(\frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \mathbf{z}_u^\top \mathbf{z}_n}\right) \approx \log(\sigma(\mathbf{z}_u^\top \mathbf{z}_v)) - \sum_{i=1}^k \log(\sigma(\mathbf{z}_u^\top \mathbf{z}_{n_i})), n_i \sim P_V$$

σ : sigmoid function, makes each term a “probability” between 0 and 1

P_V : random distribution over all nodes

Negative Sampling

$$\log\left(\frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^\top \mathbf{z}_n)}\right) \approx \log(\sigma(\mathbf{z}_u^\top \mathbf{z}_v)) - \sum_{i=1}^k \log(\sigma(\mathbf{z}_u^\top \mathbf{z}_{n_i})), n_i \sim P_V$$

- Sample k negative nodes proportionally to degree
- Two considerations for k (number of negative samples):
 1. Higher k gives more robust estimates
 2. Higher k corresponds to higher prior on negative events
In practice $k = 5..20$

Random Walks: overview

1. Run short fixed-length random walks starting from each node on the graph using some strategy R
2. For each node u collect $N_R(u)$, the multiset of nodes visited on random walks starting from u
3. Optimize embeddings using Stochastic Gradient Descent (approximating through negative sampling)

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$

Walk this (or that) way...


- We have described how to optimize embeddings given random walk statistics
- What strategies should we use to run these random walks?
 - DeepWalk (Perozzi et al., 2013): Fixed-length, unbiased random walks starting from each node
 - This notion of similarity is too constrained
- How can we generalize this?

node2vec: Overview


- Goal: Embed nodes with similar network neighborhoods close in the feature space
 - Framed as a maximum likelihood optimization problem, independent of the downstream prediction task
- Key observation: Flexible notion of network neighborhood $N_R(u)$ of node u leads to rich node embeddings
- node2vec: create biased second order random walk R to generate network neighborhood $N_R(u)$ of node u

node2vec: biased walks

- **Idea:** use flexible, biased random walks that can trade off between local and global views of the network



node2vec: biased walks



- Walk of length 3 (or neighborhood $N_R(u)$ of size 3):

$N_{BFS}(u) = \{s_1, s_2, s_3\}$ – Local microscopic view

$N_{DFS}(u) = \{s_4, s_5, s_6\}$ – Global macroscopic view

node2vec: interpolating BFS and DFS

Biased fixed-length random walk R that given a node u generates neighborhood $N_R(u)$


- Two parameters:
 - Return parameter p
Return back to the previous node
 - In-out parameter q
Moving outwards (DFS) vs. inwards (BFS)
Intuitively, q is the “ratio” of BFS vs. DFS

node2vec: biased random walks

Biased 2nd-order random walks explore network neighborhoods


Random walk just traversed edge (s_1, w) and is now at w

Neighbors of w can only be



node2vec: biased random walks


Random walk just traversed edge (s_1, w) and is now at w
Where to go next?



- p, q : model transition probabilities
 - p return parameter
 - q "walk away" parameter
- Note: $1/p, 1/q, 1$, are unnormalized probabilities

node2vec: biased random walks

Random walk just traversed edge (s_1, w) and is now at w
Where to go next?



$w \rightarrow$

Target t	Prob.	Dist. (s_1, t)
s_1	$1/p$	0
s_2	1	1
s_3	$1/q$	2
s_4	$1/q$	2

- **BFS-like** walk: low value of p
- **DFS-like** walk: low value of q

node2vec algorithm

1. Compute random walk probabilities
2. Simulate r random walks of length l starting at each node u
3. Optimize the node2vec objective using Stochastic Gradient Descent

Linear-time complexity

All 3 steps are individually parallelizable

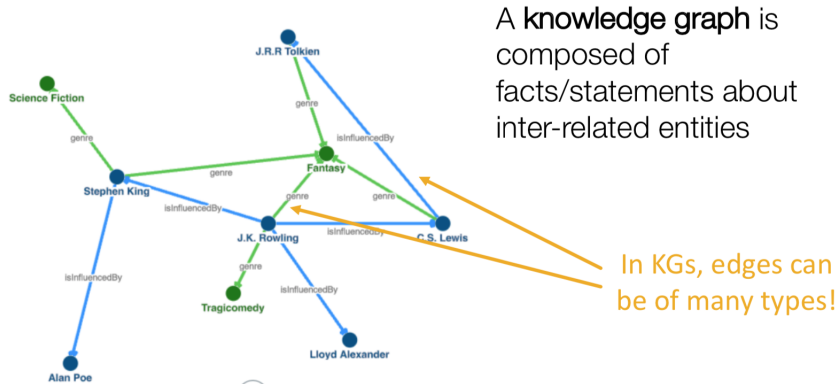
How to use the node embeddings

- **Clustering/community detection:** Cluster points z_i
- **Node classification:** Predict label $f(z_i)$ of node based on z_i
- **Link prediction:** Predict edge (i, j) based on $f(z_i, z_j)$
 - Where we can apply:
 - Concatenation: $f(z_i, z_j) = g([z_i, z_j])$
 - Hadamard product: $f(z_i, z_j) = g(z_i \odot z_j)$
 - Sum/Average: $f(z_i, z_j) = g([z_i, z_j])$
 - Distance: $f(z_i, z_j) = g(\|z_i - z_j\|_2)$

Summary

- **Basic idea:** Embed nodes so that distances in embedding space reflect node similarities in the original network
- Different notions of node similarity
 - Adjacency-based (i.e., similar if connected)
 - Multi-hop similarity definitions
 - Random walk approaches (covered today)
- No single method wins in all cases...
 - e.g. node2vec performs better on node classification while multi-hop methods performs better on link prediction
- Random walk approaches are generally more efficient
- **In general:** Must choose definition of node similarity that matches your application!

Knowledge Graph Embeddings




A **knowledge graph** is composed of facts/statements about inter-related entities

In KGs, edges can be of many types!

Nodes are referred to as **entities**,
edges as **relations**

Knowledge Graph completion




KG **incompleteness** can substantially affect the efficiency of systems relying on it!

Create a link prediction model that learns from local and global connectivity patterns in the KG, taking into account entities and relationships of different types at the same time

TransE

- In TransE, relationships between entities are represented as triplets
(head entity), (relation), (tail entity) : (h, l, t)
- Entities are first embedded in an entity space R^k
 - similarly to the previous methods
- Relations are represented as **translations**
 - $h + l \approx t$ if the fact is true
 - else, $h + l \neq t$



TransE algorithm

Algorithm 1 Learning TransE

input Training set $S = \{(h, \ell, t)\}$, entities and rel. sets E and L , margin γ , embeddings dim. k .

- 1: **initialize** $\ell \leftarrow \text{uniform}(-\frac{6}{\sqrt{k}}, \frac{6}{\sqrt{k}})$ for each $\ell \in L$
- 2: $\ell \leftarrow \ell / \|\ell\|$ for each $\ell \in L$
- 3: $e \leftarrow \text{uniform}(-\frac{6}{\sqrt{k}}, \frac{6}{\sqrt{k}})$ for each entity $e \in E$
- 4: **loop**
- 5: $e \leftarrow e / \|e\|$ for each entity $e \in E$
- 6: $S_{batch} \leftarrow \text{sample}(S, b)$ // sample a minibatch of size b
- 7: $T_{batch} \leftarrow \emptyset$ // initialize the set of pairs of triplets
- 8: **for** $(h, \ell, t) \in S_{batch}$ **do**
- 9: $(h', \ell, t') \leftarrow \text{sample}(S'_{(h, \ell, t)})$ // sample a corrupted triplet
- 10: $T_{batch} \leftarrow T_{batch} \cup \{((h, \ell, t), (h', \ell, t'))\}$
- 11: **end for**
- 12: Update embeddings w.r.t.
$$\sum_{((h, \ell, t), (h', \ell, t')) \in T_{batch}} \nabla [\gamma + \underset{\substack{\text{positive} \\ \text{sample}}}{d(\mathbf{h} + \ell, \mathbf{t})} - \underset{\substack{\text{negative} \\ \text{sample}}}{d(\mathbf{h}' + \ell, \mathbf{t}')}]_+$$
- 13: **end loop**

Entities and relations are initialized uniformly, and normalized

Negative sampling with triplet that does not appear in the KG

Comparative loss: favors lower distance values for valid triplets, high distance values for corrupted ones